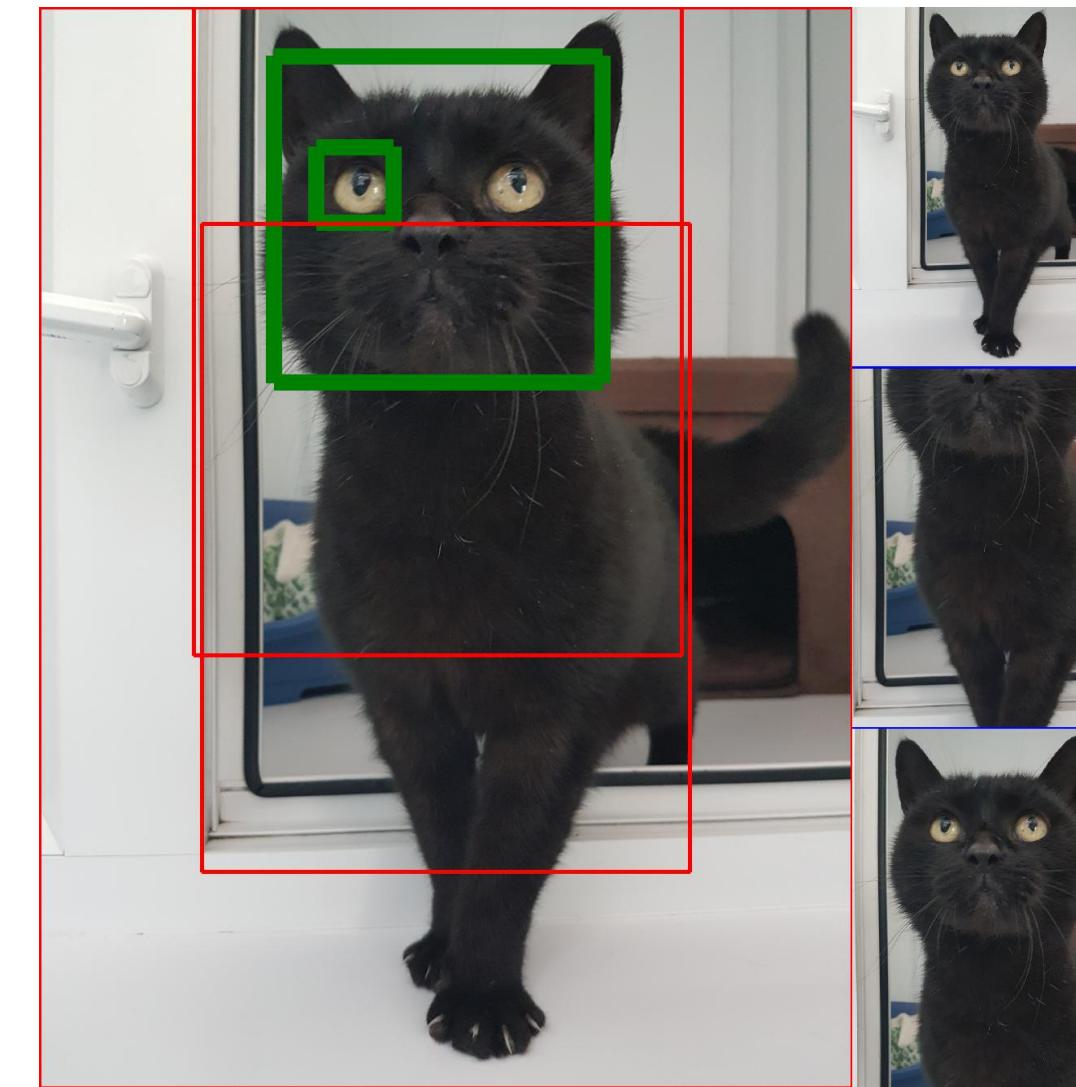


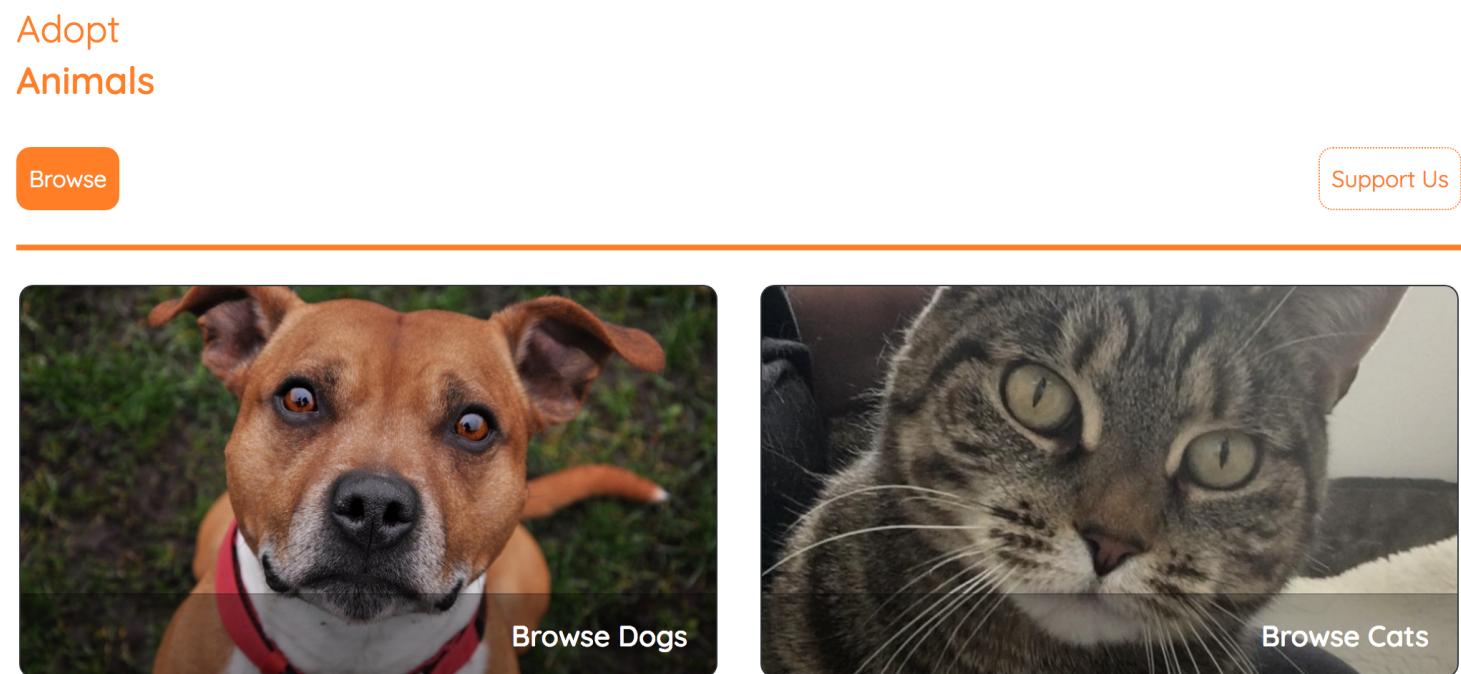
Whiskers, Ears and Puppy Dog Tails: finding out what animal pictures are made of



Who am I? Why am I doing this?

-  Mike Moran, Engineer at Skyscanner
- Adopt Animals
- Skyscanner “One Day”
- A general Machine Learning interest

Context: Adopt Animals



What is Adopt Animals?

We're building new ways to help animals in shelters find their loving forever homes.



– **Adopt Animals**, by **Kale** charity

What were we trying to do?

"Often, the images that we get are going to be slightly off-centre w.r.t. the subject, which makes the smaller previews look odd - I'll frame my thoughts as questions then statements:

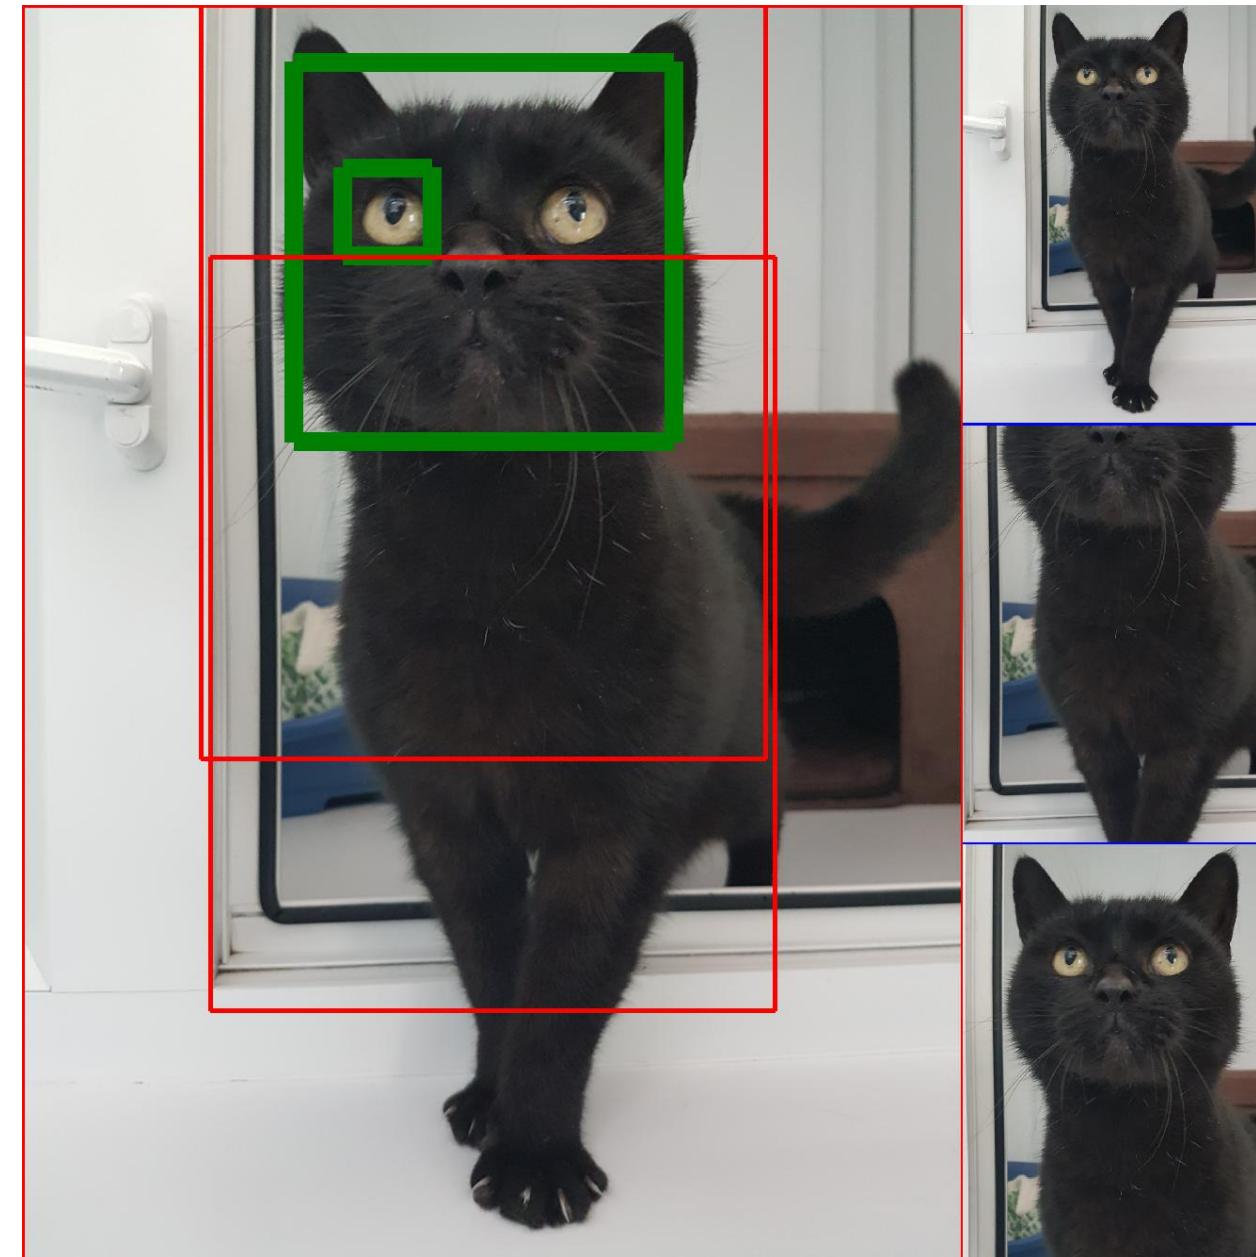
- Do ML models exist that can find the rough position of an animal's face?
- Is it per species (I'm guessing so)?
- How do I run it as part of an aggregation workflow? And what's the computational cost? Would it run on a small-ish Droplet?
- I don't need any "internal" details mapped, just where their face is, as a bounding box or something
- It doesn't have to be super perfect - more for nudging the frame in the right direction
- I can get a bunch of reference material, in the form of some snapshots of EDCH images over whatever time frame you'd like
- I can't promise to ship it immediately but it sounds like a fun side-project at least 😊"

– (Sky of Adopt Animals)

My constraints

- I ~~am lazy~~ have limited time, so an off-the-shelf solution would be good
- I am *not* a Jupyter expert (yet?)
- I have a non-uniform experience of Data-science algorithms

Ok, so how do I get to this?



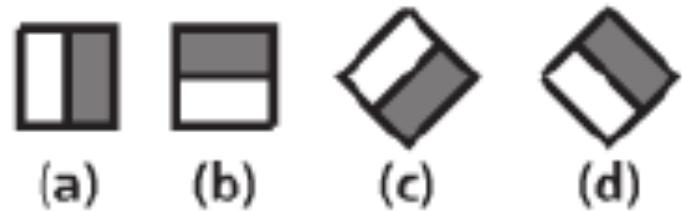


Going with something I
know: Viola Jones face
detection

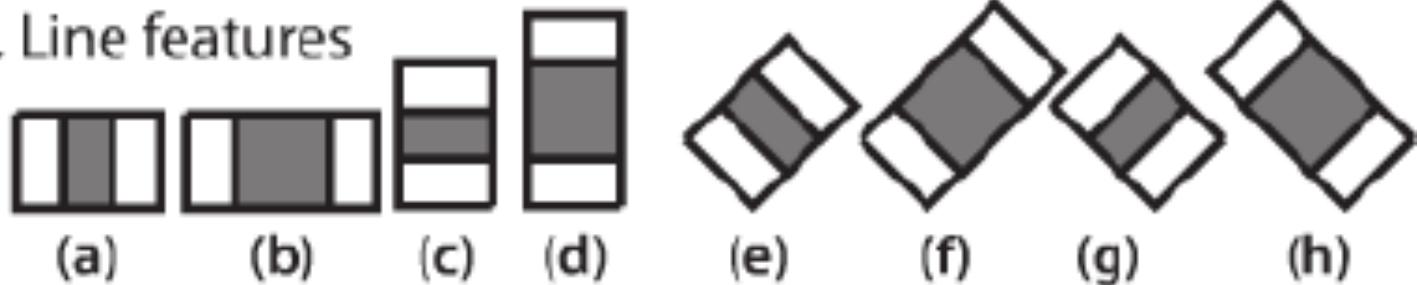
– from “Looking for Bobby but found
Paris instead”

Training: Haar features

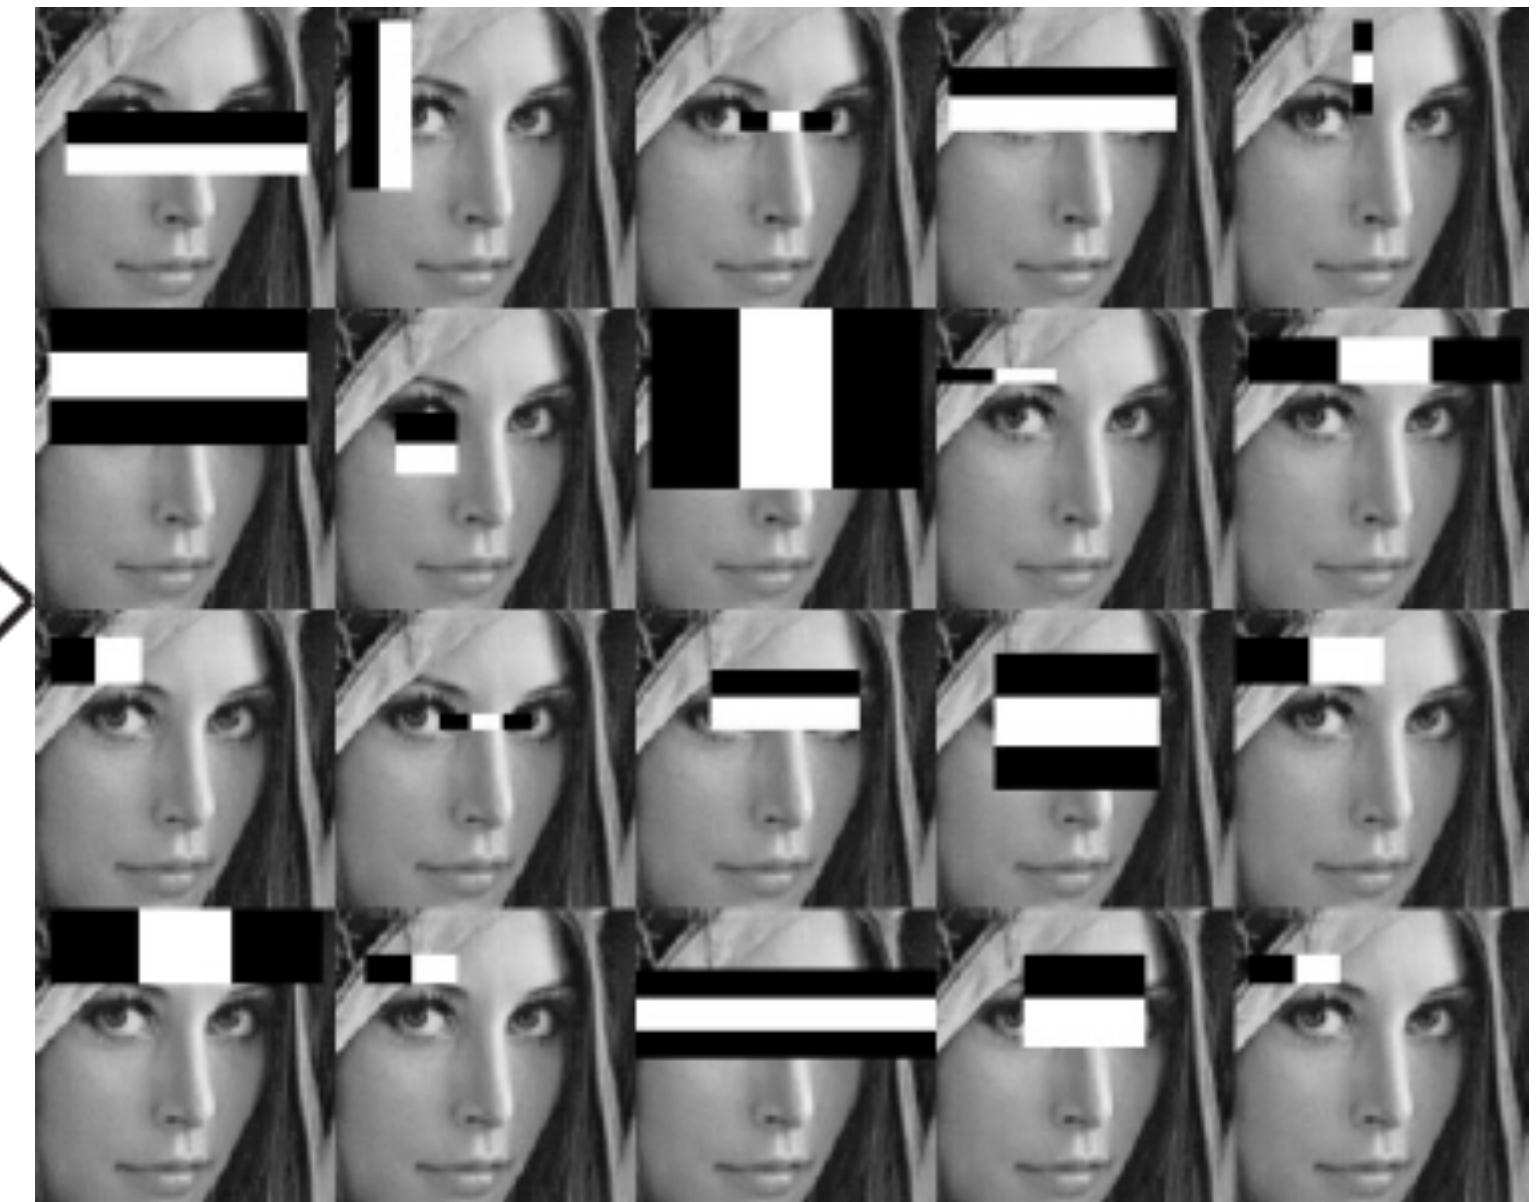
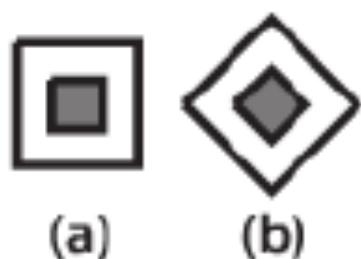
1. Edge features



2. Line features



3. Center-surround features

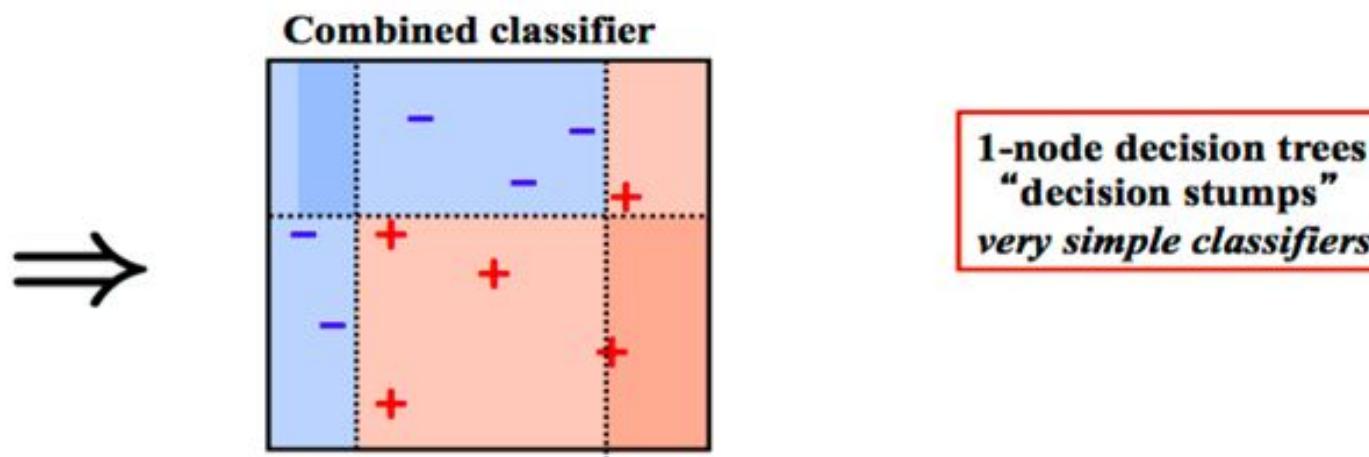


Training: Adaboost

Algorithm Adaboost - Example

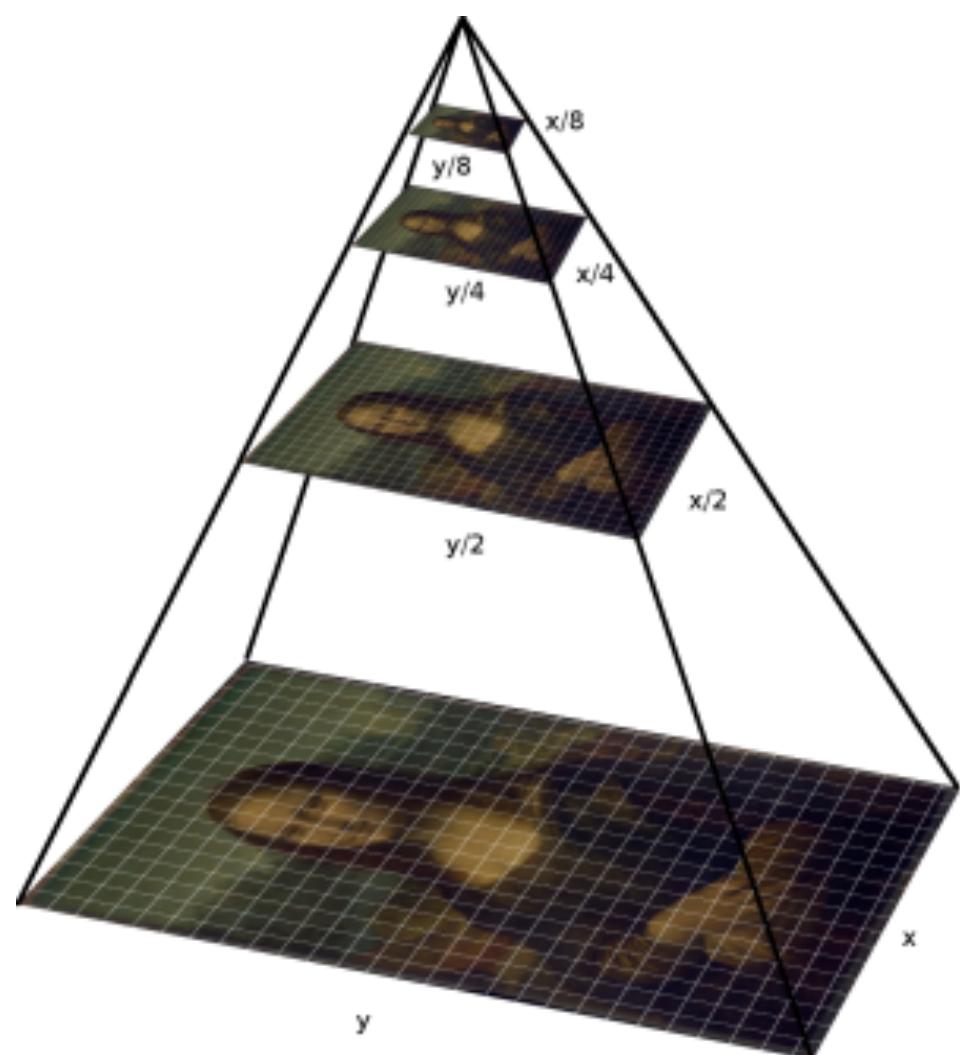
Weight each classifier and combine them:

$$.33 * \boxed{\text{blue}} + .57 * \boxed{\text{blue} \mid \text{orange}} + .42 * \boxed{\text{blue} \mid \text{orange}} \geq 0$$



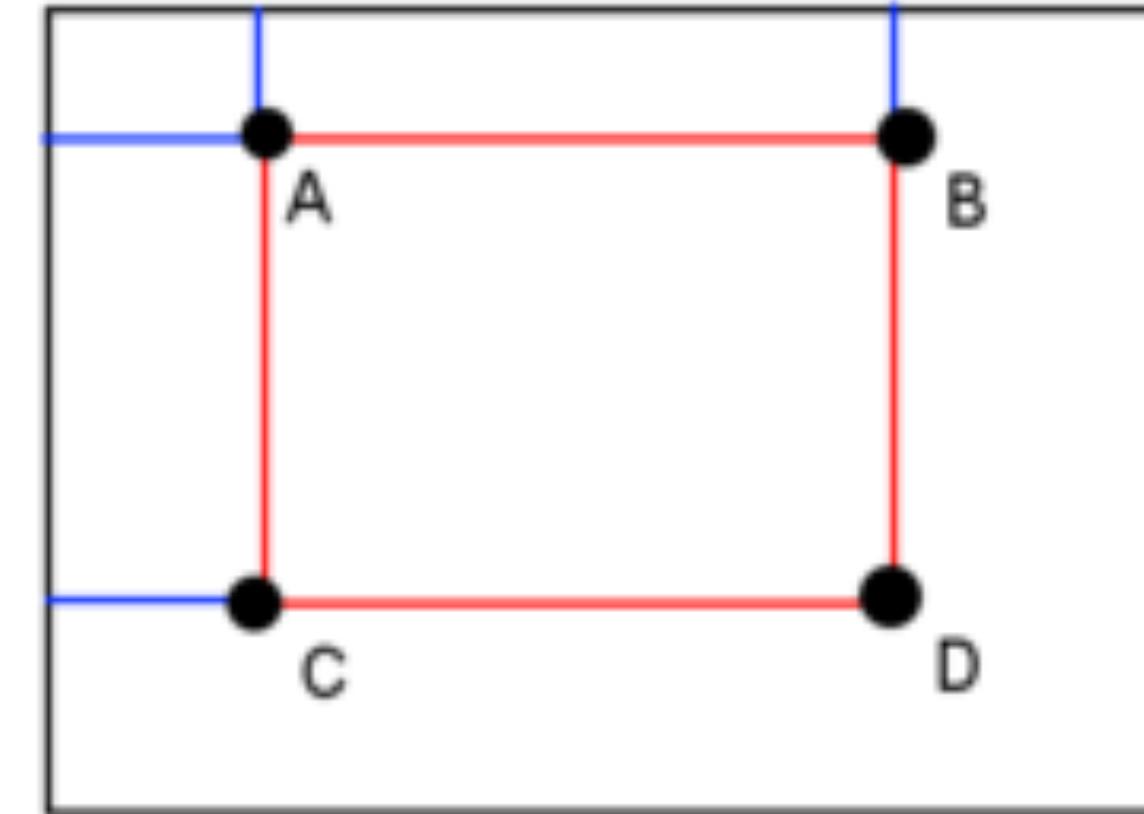
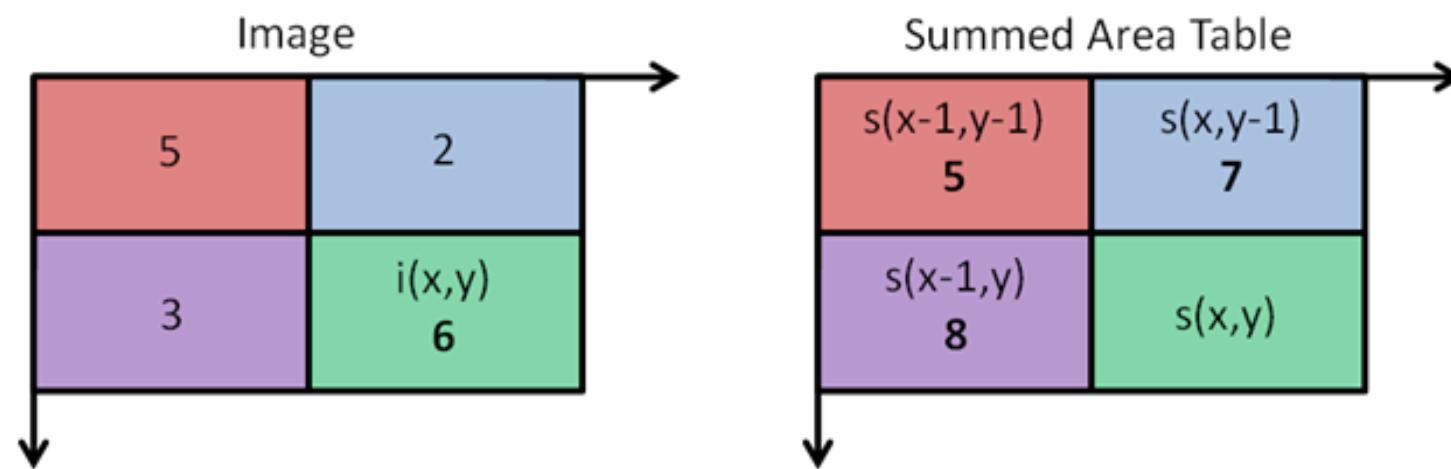
Applying: Image Pyramid

- deals with varying feature size



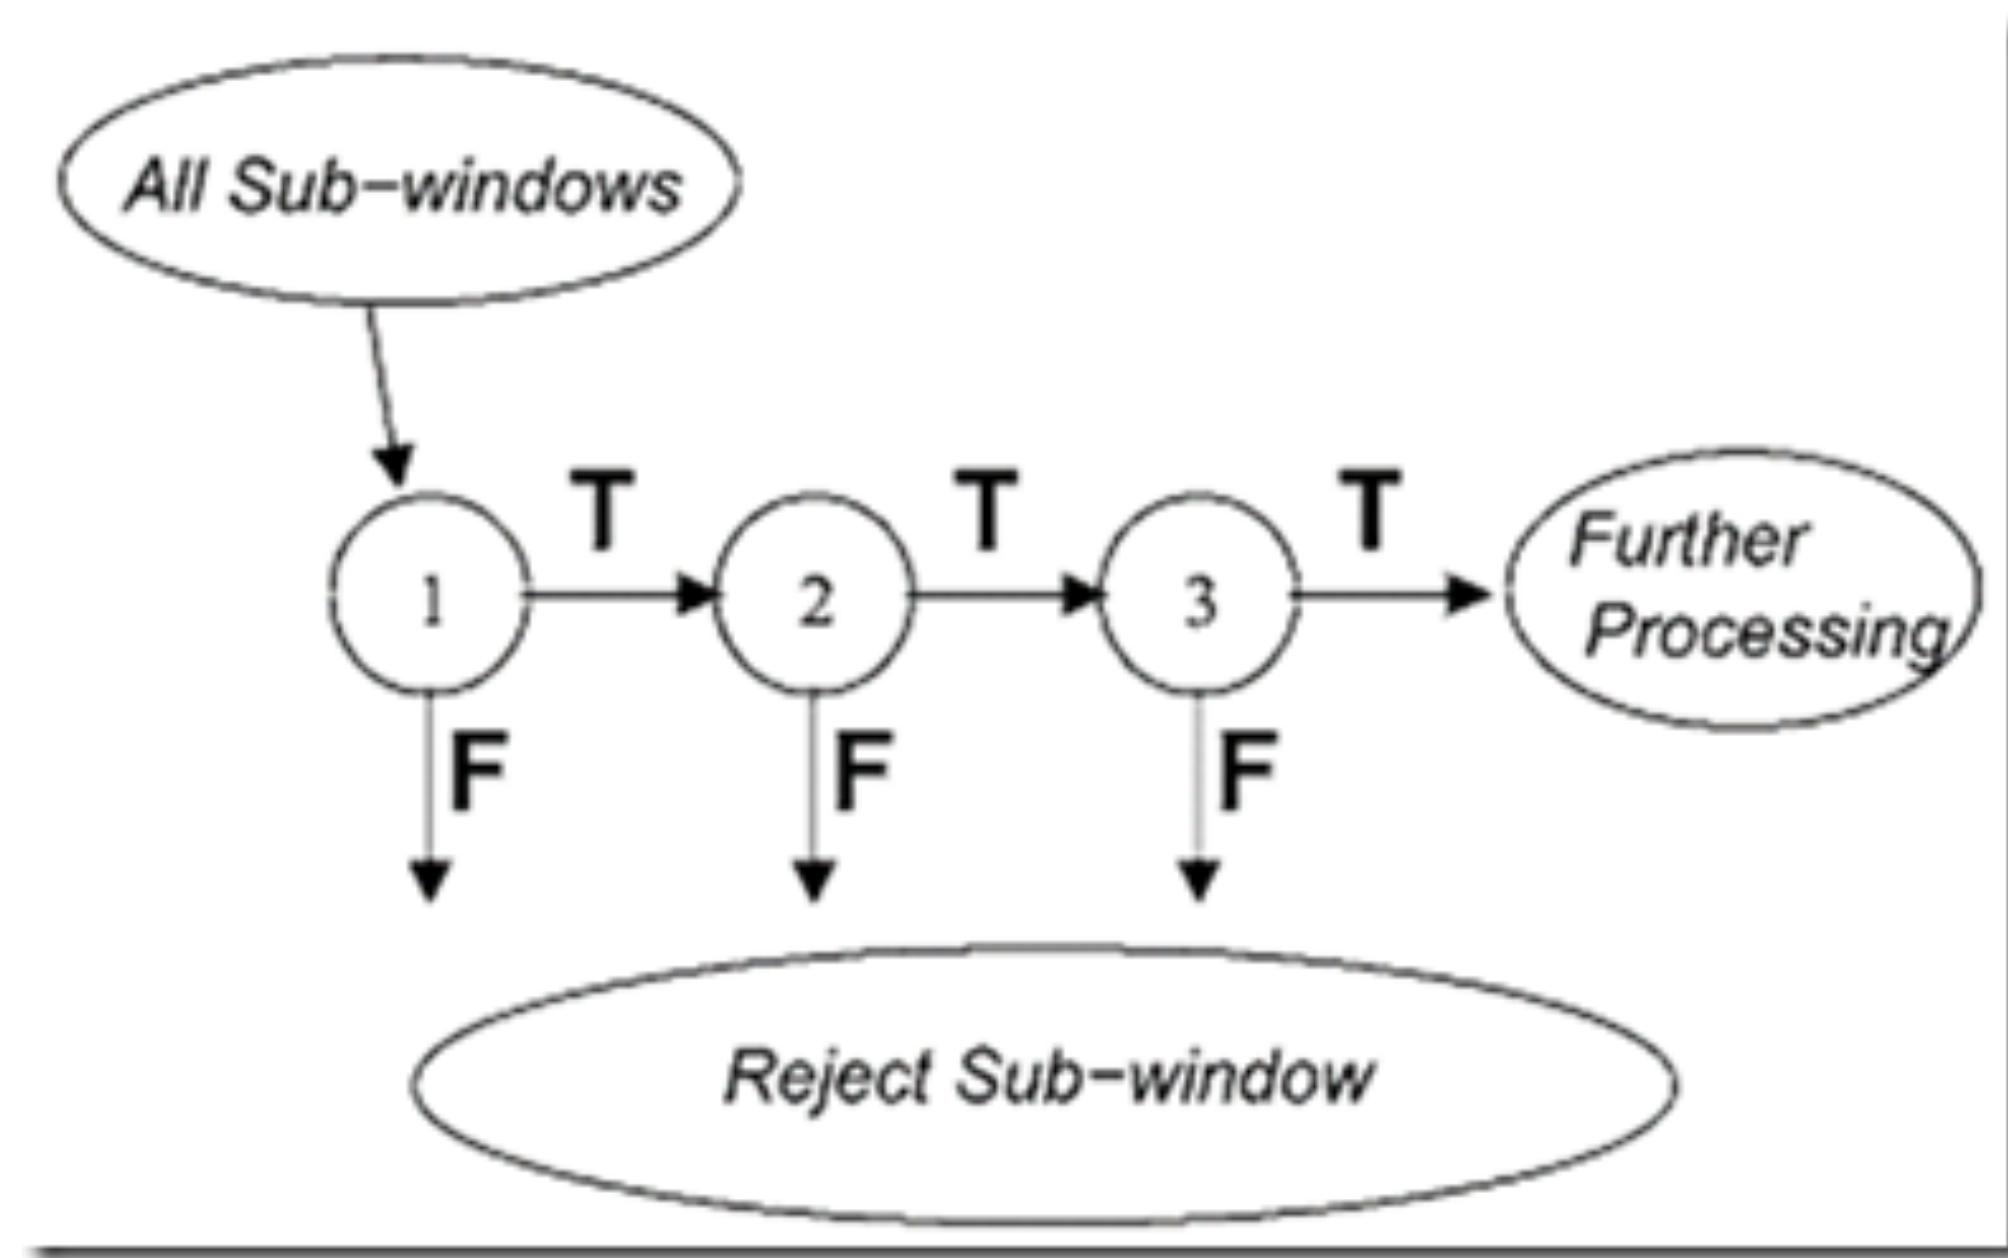
Applying: Integral Image

- efficiently evaluates hear features; $O(n^2) \rightarrow O(1)$



$$\text{Sum} = D - B - C + A$$

Applying: Haar cascade

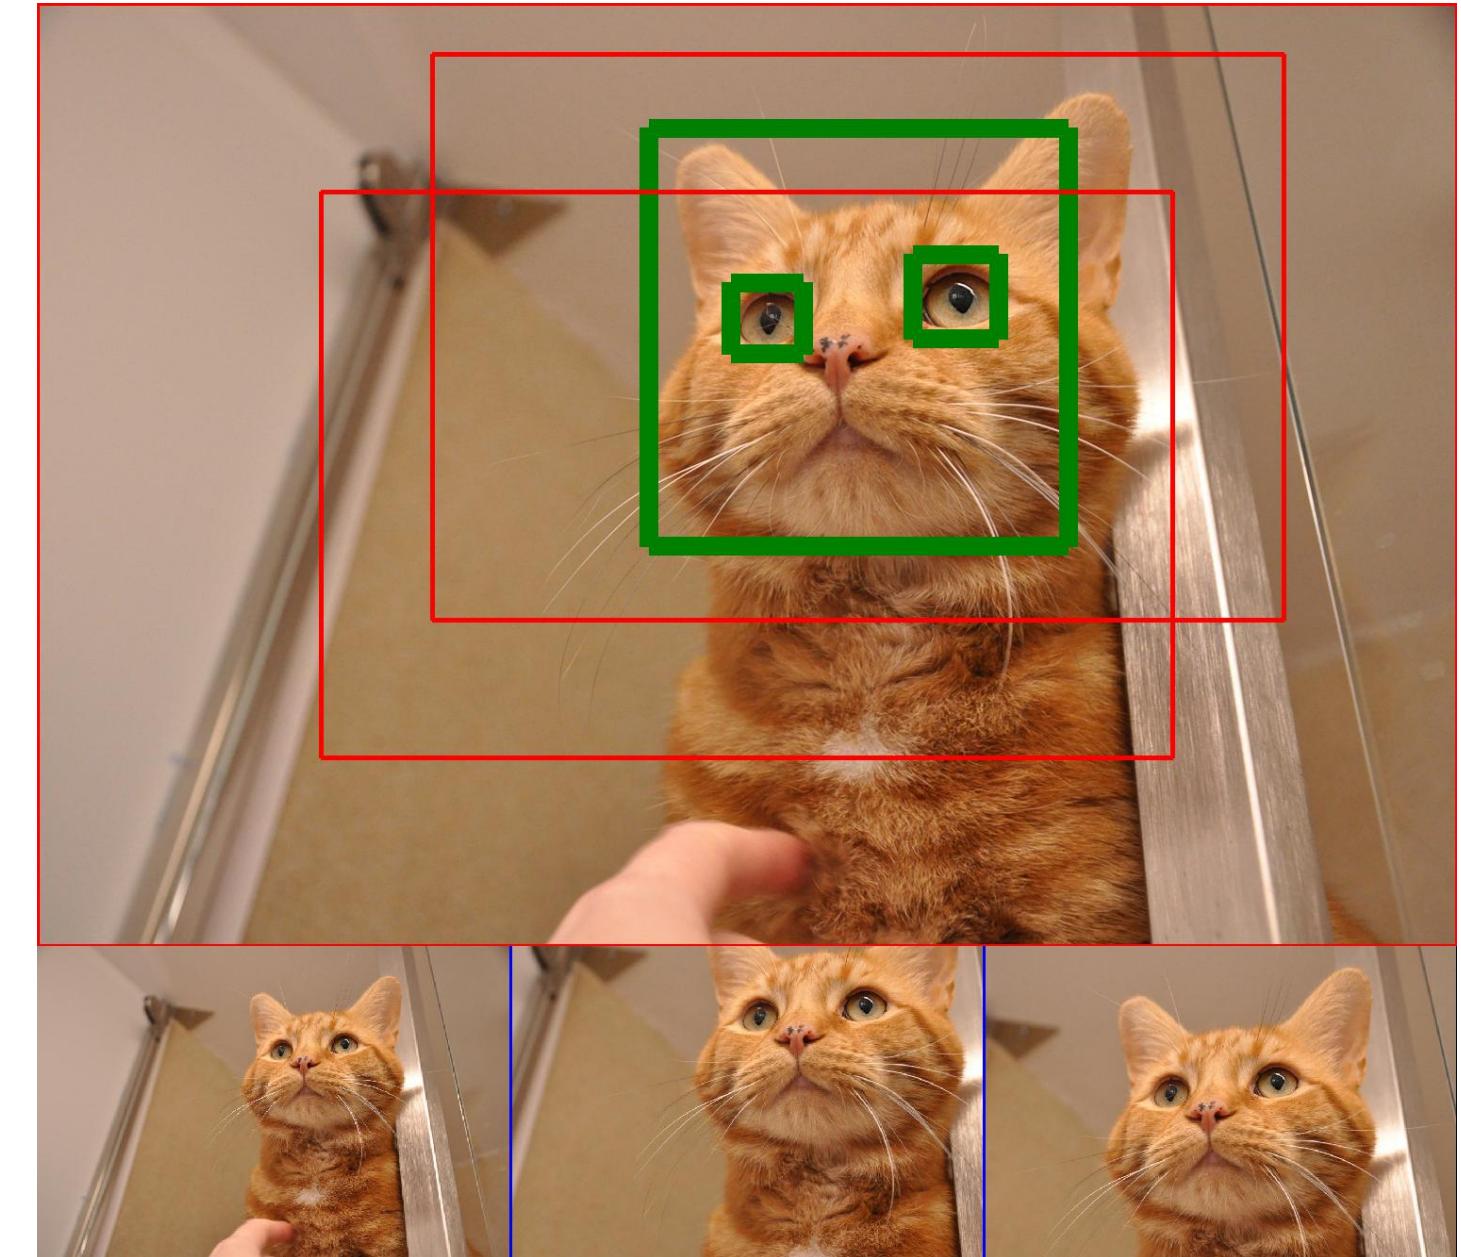
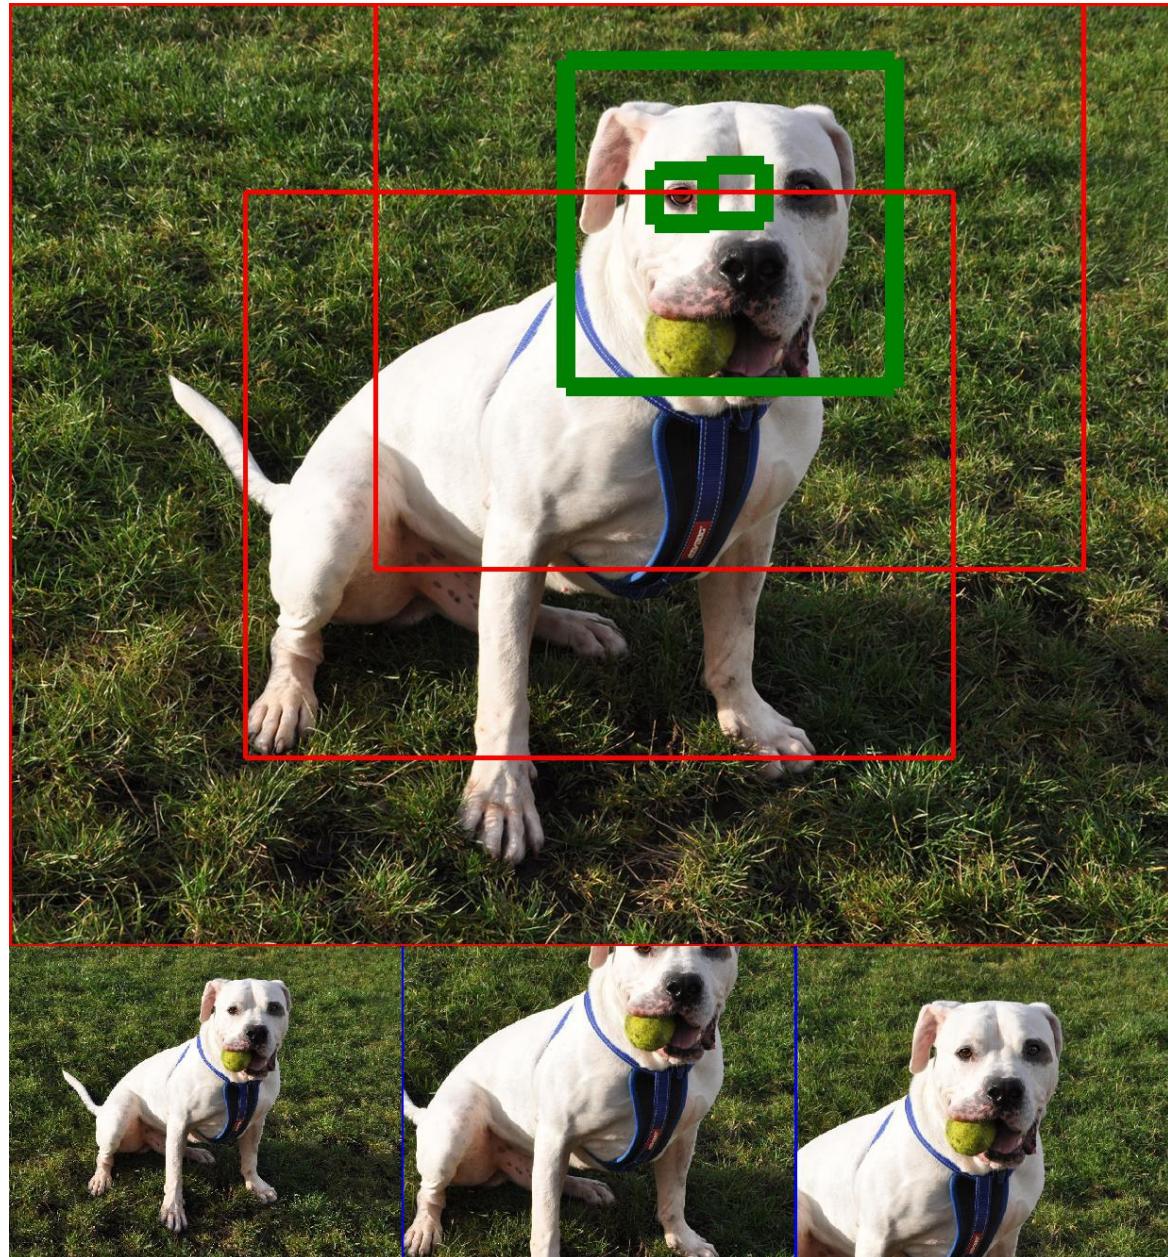


Visualising applying to image

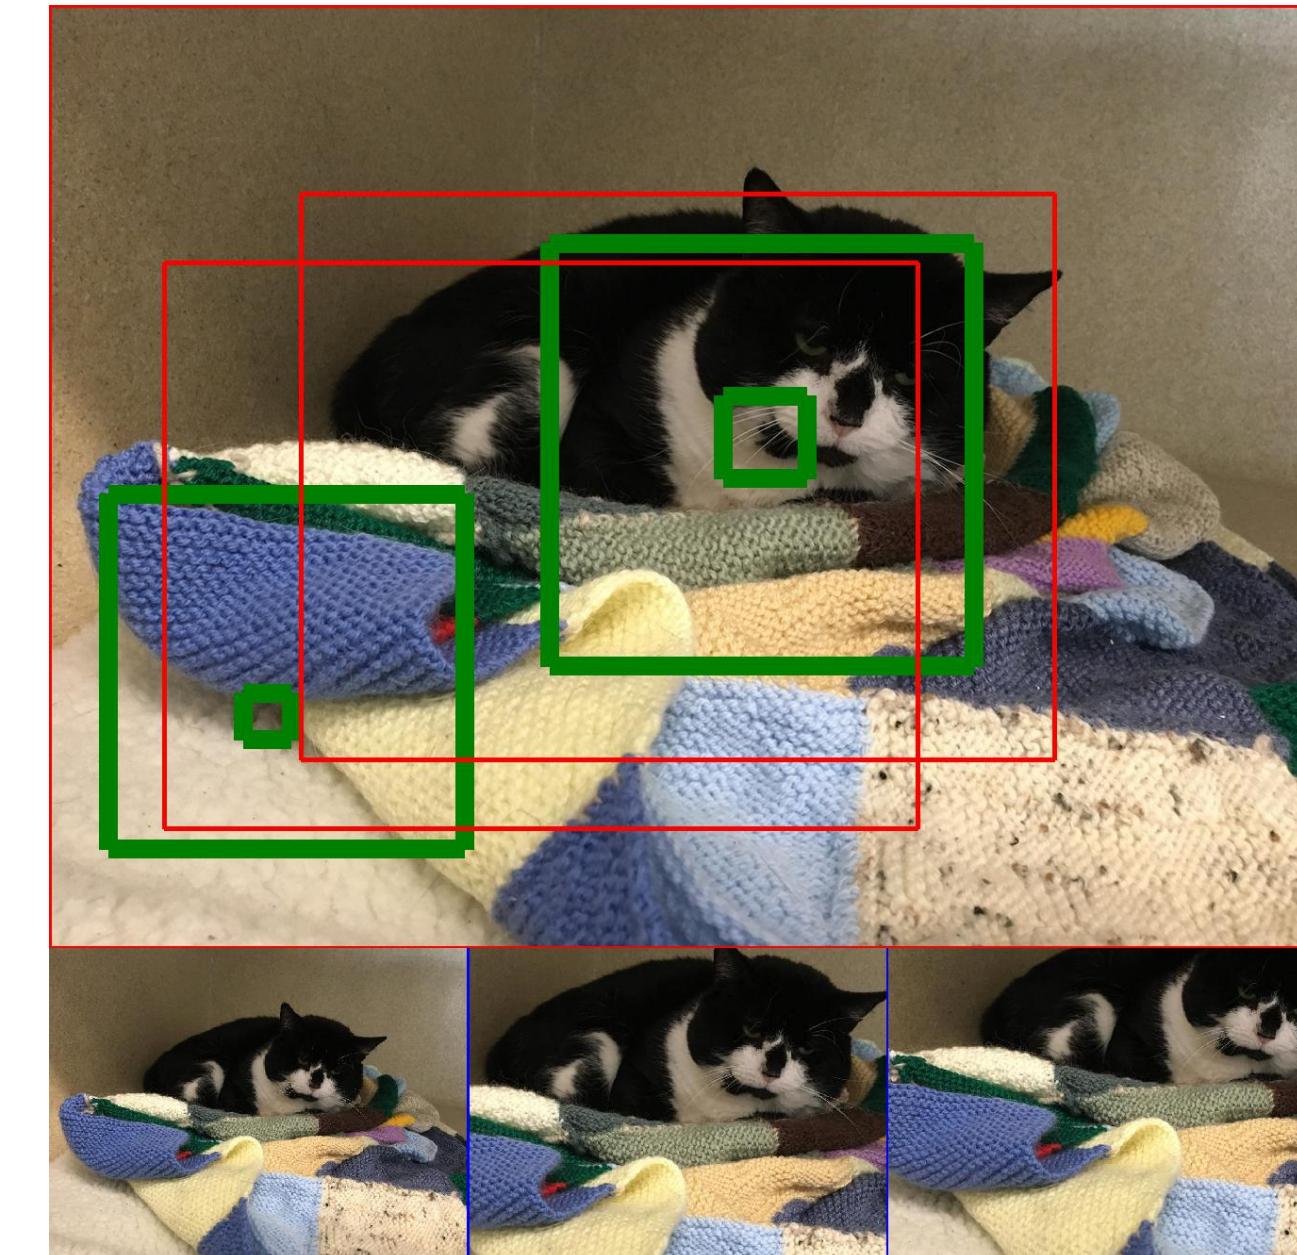
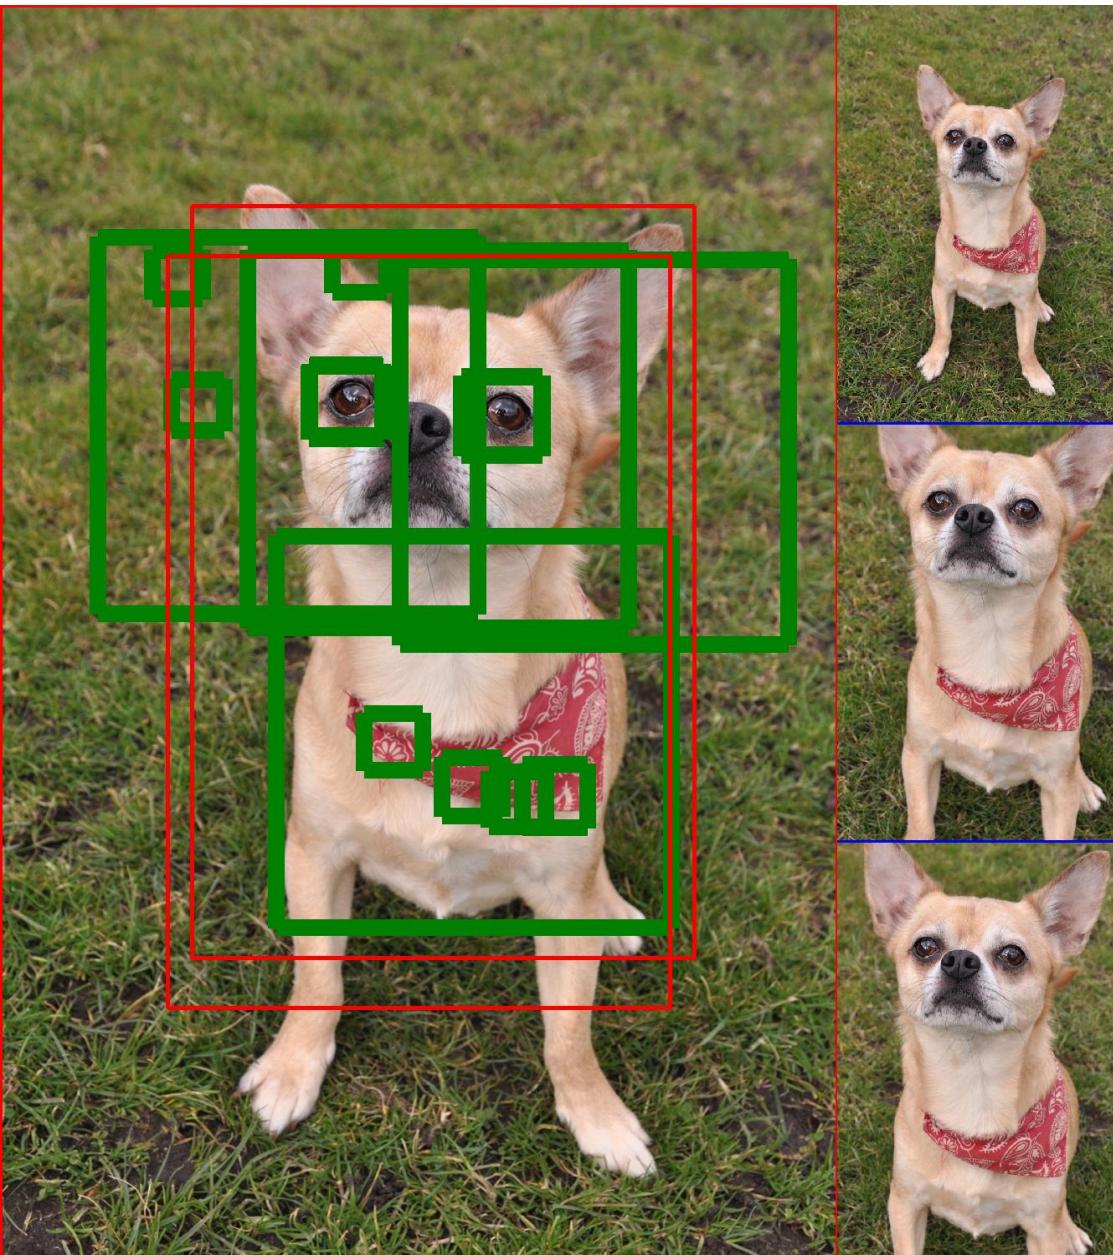


[YouTube visualisation by Adam Harvey](#)

Where have I gotten to in this? Job done, ye?



Not so fast



Visualising varying min_size

Learnings / Reinforcements

- Maybe too lazy in not focussing on understanding e.g. playing with scale-size instead of max-size parameters. Became clear once I actually spent some time on it.
- Some parameters still seem like magic (min neighbours = 5)
- You don't have to be an expert to be helpful

Backing up: what next?

- Productionising the cats-only thumbnails
- Pre-learned Dogs?
- Auto-learning of parameters on top of same algorithm
- #HASHTAG DEEP LEARNING?
 - Deep learning is expensive, both in terms of cpu cost and collecting of training data, is it worth it?

Thanks / Contact details / Questions



Sky of Adopt Animals, <https://twitter.com/carrotcodes>



Kale charity, <https://www.kale.org.uk/>



(that's me), https://twitter.com/mike_moran